

VGG-16 Model

Action Classes - 15

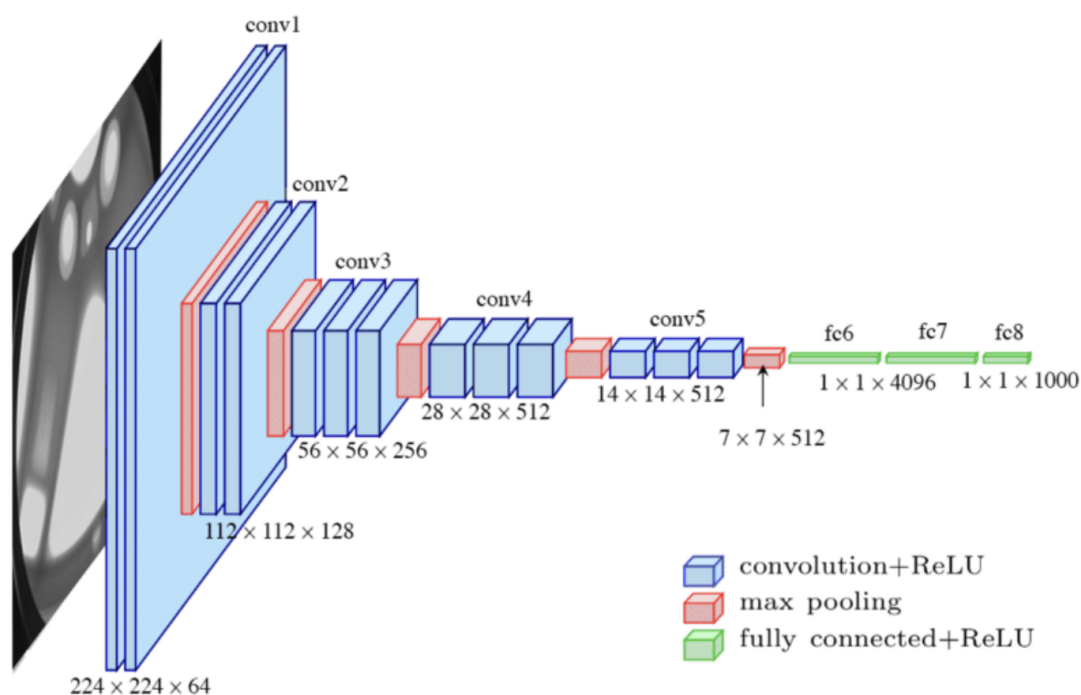
```
In [1]: from keras import models
from keras.layers import Dense, Flatten
from keras import backend as K
import numpy as np
import matplotlib.pyplot as plt

from keras.applications import vgg16
```

```
In [2]: import tensorflow as tf
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

Num GPUs Available: 1

2022-08-25 17:33:47.689083: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
 2022-08-25 17:33:47.717502: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero
 2022-08-25 17:33:47.717753: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero



Dataset

```
In [3]: from keras.preprocessing.image import ImageDataGenerator

dataset_path = "./frames/"
# will contain the categories in respective folders

# Data generators
train_datagen = ImageDataGenerator(rescale=1/255, validation_split=0.2)
```

```
In [4]: image_size = (224,224)
batch_size = 10

train_batches = train_datagen.flow_from_directory(
    dataset_path,
    target_size = image_size,
    batch_size = batch_size,
    class_mode = "categorical",
    subset = "training"
)

validation_batches = train_datagen.flow_from_directory(
    dataset_path,
    target_size = image_size,
    batch_size = batch_size,
    class_mode = "categorical",
    subset = "validation"
)

test_batches = train_datagen.flow_from_directory(
    dataset_path,
    target_size = image_size,
    batch_size = batch_size,
    class_mode = "categorical",
    subset = "validation"
)
```

Found 3767 images belonging to 15 classes.
Found 934 images belonging to 15 classes.
Found 934 images belonging to 15 classes.

```
In [5]: train_batches.class_indices
```

```
Out[5]: {'ApplyLipstick': 0,
        'Archery': 1,
        'BabyCrawling': 2,
        'Biking': 3,
        'Diving': 4,
        'Fencing': 5,
        'Kayaking': 6,
        'MilitaryParade': 7,
        'PizzaTossing': 8,
        'ShavingBeard': 9,
        'SkateBoarding': 10,
        'SumoWrestling': 11,
        'TennisSwing': 12,
        'Typing': 13,
        'WritingOnBoard': 14}
```

```
In [6]: from matplotlib import pyplot as plt

def plot_images(images_arr):
    fig, axes = plt.subplots(1,10)
    axes = axes.flatten()
    for img, ax in zip(images_arr, axes):
        ax.imshow(img)
        ax.axis('off')
    plt.tight_layout()
    plt.show()
```

```
In [7]: imgs, labels = train_batches[0]
plot_images(imgs)
print(labels[:10])
```



```
[[1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]
 [0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]
```

Initialize Model

```
In [8]: vggmodeltop = vgg16.VGG16(include_top=True,
                                   input_shape=(224,224,3),
                                   pooling='avg',
                                   weights='imagenet')

for (i,layer) in enumerate(vggmodeltop.layers):
    layer.trainable = False
    print((i, layer.name, layer.output_shape))
```

2022-08-25 17:33:48.451468: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

2022-08-25 17:33:48.451975: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-08-25 17:33:48.452250: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-08-25 17:33:48.452522: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-08-25 17:33:48.887809: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-08-25 17:33:48.887991: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-08-25 17:33:48.888114: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:975] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so returning NUMA node zero

2022-08-25 17:33:48.888215: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1532] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 3368 MB memory: -> device: 0, name: NVIDIA GeForce GTX 1050 Ti, pci bus id: 0000:01:00.0, compute capability: 6.1

(0, 'input_1', [(None, 224, 224, 3)])

(1, 'block1_conv1', (None, 224, 224, 64))

(2, 'block1_conv2', (None, 224, 224, 64))

(3, 'block1_pool', (None, 112, 112, 64))

(4, 'block2_conv1', (None, 112, 112, 128))

(5, 'block2_conv2', (None, 112, 112, 128))

(6, 'block2_pool', (None, 56, 56, 128))

(7, 'block3_conv1', (None, 56, 56, 256))

(8, 'block3_conv2', (None, 56, 56, 256))

(9, 'block3_conv3', (None, 56, 56, 256))

(10, 'block3_pool', (None, 28, 28, 256))

(11, 'block4_conv1', (None, 28, 28, 512))

(12, 'block4_conv2', (None, 28, 28, 512))

(13, 'block4_conv3', (None, 28, 28, 512))

(14, 'block4_pool', (None, 14, 14, 512))

(15, 'block5_conv1', (None, 14, 14, 512))

(16, 'block5_conv2', (None, 14, 14, 512))

(17, 'block5_conv3', (None, 14, 14, 512))

(18, 'block5_pool', (None, 7, 7, 512))

(19, 'flatten', (None, 25088))

(20, 'fc1', (None, 4096))

(21, 'fc2', (None, 4096))

(22, 'predictions', (None, 1000))

```
In [9]: vggmodel = vgg16.VGG16(include_top=False,
                                input_shape=(224,224,3),
                                pooling='avg',classes=15,
                                weights='imagenet')

for (i,layer) in enumerate(vggmodel.layers):
    layer.trainable = False
    print((i, layer.name, layer.output_shape))
```

```

(0, 'input_2', [(None, 224, 224, 3)])
(1, 'block1_conv1', (None, 224, 224, 64))
(2, 'block1_conv2', (None, 224, 224, 64))
(3, 'block1_pool', (None, 112, 112, 64))
(4, 'block2_conv1', (None, 112, 112, 128))
(5, 'block2_conv2', (None, 112, 112, 128))
(6, 'block2_pool', (None, 56, 56, 128))
(7, 'block3_conv1', (None, 56, 56, 256))
(8, 'block3_conv2', (None, 56, 56, 256))
(9, 'block3_conv3', (None, 56, 56, 256))
(10, 'block3_pool', (None, 28, 28, 256))
(11, 'block4_conv1', (None, 28, 28, 512))
(12, 'block4_conv2', (None, 28, 28, 512))
(13, 'block4_conv3', (None, 28, 28, 512))
(14, 'block4_pool', (None, 14, 14, 512))
(15, 'block5_conv1', (None, 14, 14, 512))
(16, 'block5_conv2', (None, 14, 14, 512))
(17, 'block5_conv3', (None, 14, 14, 512))
(18, 'block5_pool', (None, 7, 7, 512))
(19, 'global_average_pooling2d', (None, 512))

```

```

In [10]: model = models.Sequential()

dense_layer_1 = Dense(32, activation='relu')
dense_layer_2 = Dense(32, activation='relu')
prediction_layer = Dense(15, activation='softmax')

model.add(vggmodel)
model.add(dense_layer_1)
model.add(dense_layer_2)
model.add(prediction_layer)

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Functional)	(None, 512)	14714688
dense (Dense)	(None, 32)	16416
dense_1 (Dense)	(None, 32)	1056
dense_2 (Dense)	(None, 15)	495
=====		
Total params: 14,732,655		
Trainable params: 17,967		
Non-trainable params: 14,714,688		

```

In [11]: model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy'],
)

```

```

In [12]: model.save("./models/action-class-15-model-vgg16.h5")

```

```

In [13]: fit = model.fit(train_batches, epochs=20, validation_data=validation_batches)

Epoch 1/20

```

```
2022-08-25 17:33:52.383352: I tensorflow/stream_executor/cuda/cuda_dnn.c
c:384] Loaded cuDNN version 8401
2022-08-25 17:33:52.797291: I tensorflow/core/platform/default/subproces
s.cc:304] Start cannot spawn child process: No such file or directory
2022-08-25 17:33:53.014079: W tensorflow/core/common_runtime/bfc_allocato
r.cc:290] Allocator (GPU_0_bfc) ran out of memory trying to allocate 2.35
GiB with freed_by_count=0. The caller indicates that this is not a failur
e, but this may mean that there could be performance gains if more memory
were available.
```

```
377/377 [=====] - 71s 180ms/step - loss: 2.0605
- accuracy: 0.3770 - val_loss: 1.4728 - val_accuracy: 0.5728
Epoch 2/20
377/377 [=====] - 67s 177ms/step - loss: 1.0353
- accuracy: 0.7075 - val_loss: 0.9628 - val_accuracy: 0.6938
Epoch 3/20
377/377 [=====] - 69s 183ms/step - loss: 0.6941
- accuracy: 0.8025 - val_loss: 0.8382 - val_accuracy: 0.7152
Epoch 4/20
377/377 [=====] - 79s 209ms/step - loss: 0.5286
- accuracy: 0.8439 - val_loss: 0.7389 - val_accuracy: 0.7484
Epoch 5/20
377/377 [=====] - 108s 286ms/step - loss: 0.4212
- accuracy: 0.8832 - val_loss: 0.6552 - val_accuracy: 0.7741
Epoch 6/20
377/377 [=====] - 110s 291ms/step - loss: 0.3491
- accuracy: 0.9005 - val_loss: 0.6646 - val_accuracy: 0.7655
Epoch 7/20
377/377 [=====] - 111s 293ms/step - loss: 0.2894
- accuracy: 0.9196 - val_loss: 0.6429 - val_accuracy: 0.7687
Epoch 8/20
377/377 [=====] - 110s 290ms/step - loss: 0.2431
- accuracy: 0.9326 - val_loss: 0.6903 - val_accuracy: 0.7602
Epoch 9/20
377/377 [=====] - 111s 294ms/step - loss: 0.2108
- accuracy: 0.9403 - val_loss: 0.5873 - val_accuracy: 0.7944
Epoch 10/20
377/377 [=====] - 112s 297ms/step - loss: 0.1868
- accuracy: 0.9448 - val_loss: 0.5973 - val_accuracy: 0.7912
Epoch 11/20
377/377 [=====] - 111s 294ms/step - loss: 0.1578
- accuracy: 0.9586 - val_loss: 0.6983 - val_accuracy: 0.7623
Epoch 12/20
377/377 [=====] - 112s 296ms/step - loss: 0.1384
- accuracy: 0.9610 - val_loss: 0.8109 - val_accuracy: 0.7281
Epoch 13/20
377/377 [=====] - 111s 293ms/step - loss: 0.1313
- accuracy: 0.9631 - val_loss: 0.6609 - val_accuracy: 0.7762
Epoch 14/20
377/377 [=====] - 114s 302ms/step - loss: 0.1123
- accuracy: 0.9711 - val_loss: 0.6520 - val_accuracy: 0.7827
Epoch 15/20
377/377 [=====] - 112s 297ms/step - loss: 0.0956
- accuracy: 0.9774 - val_loss: 0.6752 - val_accuracy: 0.7827
Epoch 16/20
377/377 [=====] - 113s 297ms/step - loss: 0.0930
- accuracy: 0.9764 - val_loss: 0.7154 - val_accuracy: 0.7762
Epoch 17/20
377/377 [=====] - 112s 297ms/step - loss: 0.0788
- accuracy: 0.9814 - val_loss: 0.7033 - val_accuracy: 0.7794
Epoch 18/20
377/377 [=====] - 112s 297ms/step - loss: 0.0702
- accuracy: 0.9814 - val_loss: 0.6772 - val_accuracy: 0.7848
Epoch 19/20
377/377 [=====] - 113s 298ms/step - loss: 0.0643
- accuracy: 0.9851 - val_loss: 0.7302 - val_accuracy: 0.7859
Epoch 20/20
377/377 [=====] - 112s 296ms/step - loss: 0.0577
- accuracy: 0.9859 - val_loss: 0.7463 - val_accuracy: 0.7837
```

```
In [14]: model.save("./models/action-class-15-trained-vgg16.h5")
```

Evaluate and Predict

```
In [15]: model = models.load_model("./models/action-class-15-trained-vgg16.h5")
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====	=====	=====
vgg16 (Functional)	(None, 512)	14714688
dense (Dense)	(None, 32)	16416
dense_1 (Dense)	(None, 32)	1056
dense_2 (Dense)	(None, 15)	495
=====	=====	=====
Total params: 14,732,655		
Trainable params: 17,967		
Non-trainable params: 14,714,688		

```
In [16]: model.evaluate(test_batches)
```

94/94 [=====] - 26s 233ms/step - loss: 0.7463 - accuracy: 0.7837

```
Out[16]: [0.746271014213562, 0.783725917339325]
```

```
In [17]: plt.style.use("ggplot")
plt.figure()
```

```
plt.plot(np.arange(0, 20), fit.history["loss"], label="train_loss")
plt.plot(np.arange(0, 20), fit.history["val_loss"], label="val_loss")
plt.title("Training Loss")
plt.xlabel("Epoch #")
plt.ylabel("Loss")
plt.legend(loc="lower left")
plt.show()
```

```
plt.plot(np.arange(0, 20), fit.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, 20), fit.history["val_accuracy"], label="val_acc")
plt.title("Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Accuracy")
plt.legend(loc="lower left")
plt.show()
```